

Homework8

Patrick Seebold

This homework will explore basics of predictive modeling using an online-available data set concerning bike data in Seoul, South Korea.

First, let's get our libraries set:

```
library(tidymodels)
```

```
Warning: package 'tidymodels' was built under R version 4.3.3
```

```
-- Attaching packages ----- tidymodels 1.2.0 --
```

v broom	1.0.5	v recipes	1.1.0
v dials	1.3.0	v rsample	1.2.1
v dplyr	1.1.4	v tibble	3.2.1
v ggplot2	3.5.1	v tidyr	1.3.1
v infer	1.0.7	v tune	1.2.1
v modeldata	1.4.0	v workflows	1.1.4
v parsnip	1.2.1	v workflowsets	1.1.0
v purrr	1.0.2	v yardstick	1.3.1

```
Warning: package 'dials' was built under R version 4.3.3
```

```
Warning: package 'scales' was built under R version 4.3.3
```

```
Warning: package 'dplyr' was built under R version 4.3.3
```

```
Warning: package 'ggplot2' was built under R version 4.3.3
```

```
Warning: package 'infer' was built under R version 4.3.3
```

Warning: package 'modeldata' was built under R version 4.3.3

Warning: package 'parsnip' was built under R version 4.3.3

Warning: package 'purrr' was built under R version 4.3.3

Warning: package 'recipes' was built under R version 4.3.3

Warning: package 'rsample' was built under R version 4.3.3

Warning: package 'tidyr' was built under R version 4.3.3

Warning: package 'tune' was built under R version 4.3.3

Warning: package 'workflows' was built under R version 4.3.3

Warning: package 'workflowsets' was built under R version 4.3.3

Warning: package 'yardstick' was built under R version 4.3.3

```
-- Conflicts ----- tidymodels_conflicts() --  
x purrr::discard() masks scales::discard()  
x dplyr::filter()  masks stats::filter()  
x dplyr::lag()     masks stats::lag()  
x recipes::step()  masks stats::step()  
* Dig deeper into tidy modeling with R at https://www.tmw.r.org
```

```
library(lubridate)
```

Attaching package: 'lubridate'

The following objects are masked from 'package:base':

date, intersect, setdiff, union

```
library(dplyr)
library(ggplot2)
```

EDA

Next, let's load the data in and resolve the known error mentioned in the instructions. We can specify the fileEncoding and avoid checking names to solve the problem:

```
data = read.csv("SeoulBikeData.csv", fileEncoding = "Latin1", check.names = F)
```

Great, now we can go ahead with our EDA to make sure everything looks good. First, let's check for missing values:

```
sum(is.na(data))
```

```
[1] 0
```

Sweet, no missing values. makes for a good start. Next, we'll take a look at the column types, rename to avoid strange variable names, and change our categorical variables to factors. The instructions say to change our variable names later in the process, but I'd rather they be in a clean format for the rest of the EDA!

```
summary(data)
```

Date	Rented Bike Count	Hour	Temperature(°C)
Length:8760	Min. : 0.0	Min. : 0.00	Min. : -17.80
Class :character	1st Qu.: 191.0	1st Qu.: 5.75	1st Qu.: 3.50
Mode :character	Median : 504.5	Median : 11.50	Median : 13.70
	Mean : 704.6	Mean : 11.50	Mean : 12.88
	3rd Qu.: 1065.2	3rd Qu.: 17.25	3rd Qu.: 22.50
	Max. : 3556.0	Max. : 23.00	Max. : 39.40
Humidity(%)	Wind speed (m/s)	Visibility (10m)	Dew point temperature(°C)
Min. : 0.00	Min. : 0.000	Min. : 27	Min. : -30.600
1st Qu.: 42.00	1st Qu.: 0.900	1st Qu.: 940	1st Qu.: -4.700
Median : 57.00	Median : 1.500	Median : 1698	Median : 5.100
Mean : 58.23	Mean : 1.725	Mean : 1437	Mean : 4.074
3rd Qu.: 74.00	3rd Qu.: 2.300	3rd Qu.: 2000	3rd Qu.: 14.800
Max. : 98.00	Max. : 7.400	Max. : 2000	Max. : 27.200
Solar Radiation (MJ/m2)	Rainfall(mm)	Snowfall (cm)	Seasons

Min. :0.0000	Min. : 0.0000	Min. :0.00000	Length:8760
1st Qu.:0.0000	1st Qu.: 0.0000	1st Qu.:0.00000	Class :character
Median :0.0100	Median : 0.0000	Median :0.00000	Mode :character
Mean :0.5691	Mean : 0.1487	Mean :0.07507	
3rd Qu.:0.9300	3rd Qu.: 0.0000	3rd Qu.:0.00000	
Max. :3.5200	Max. :35.0000	Max. :8.80000	
Holiday	Functioning Day		
Length:8760	Length:8760		
Class :character	Class :character		
Mode :character	Mode :character		

```
colnames(data) = c('Date', 'NumBikes', "Hour", "Temperature", "Humidity",
                  "WindSpeed", "Visibility", "DewPointTemp", "SolarRad", "Rainfall",
                  "Snowfall", "Season", "Holiday", "FunctioningDay" )
```

```
data$Season = as.factor(data$Season)
data$Holiday = as.factor(data$Holiday)
data$FunctioningDay = as.factor(data$FunctioningDay)
levels(data$Season)
```

```
[1] "Autumn" "Spring" "Summer" "Winter"
```

```
levels(data$Holiday)
```

```
[1] "Holiday" "No Holiday"
```

```
levels(data$FunctioningDay)
```

```
[1] "No" "Yes"
```

Great, now we can see that everything looks good with our numerical variables, our variable names are no longer likely to cause an issue, and our categorical variables are appropriately recast as factors! Next, we will change the date into a workable arithmetic form using lubridate:

```
typeof(data$Date) # currently character
```

```
[1] "character"
```

```
data$Date = lubridate::dmy(data$Date)
typeof(data$Date) # now it's a double!
```

```
[1] "double"
```

Now that we have cleaned the data up, we will do some summary stats, specifically looking at bike rental count, rainfall, and snowfall. We'll also examine these across some categorical variables, such as FunctioningDay, Holiday, and Season:

```
data |> # check bike numbers by season
  group_by(FunctioningDay) |>
  summarize(mean = mean(NumBikes), sd = sd(NumBikes)) # Non-functioning days mean no bikes!
```

```
# A tibble: 2 x 3
  FunctioningDay mean    sd
  <fct>          <dbl> <dbl>
1 No              0      0
2 Yes            729.   642.
```

```
sum(data$FunctioningDay == "No") # 295 of the data points can be excluded
```

```
[1] 295
```

```
sub_data = subset(data, data$FunctioningDay == "Yes") # we can ignore days that are not functioning
```

```
sub_data |> # check bike numbers by season
  group_by(Season) |>
  summarize(mean = mean(NumBikes), sd = sd(NumBikes))
```

```
# A tibble: 4 x 3
  Season mean    sd
  <fct>  <dbl> <dbl>
1 Autumn  924.  618.
2 Spring  746.  619.
3 Summer 1034.  690.
4 Winter  226.  150.
```

```
sub_data |> # check bike numbers by holiday
  group_by(Holiday) |>
  summarize(mean = mean(NumBikes), sd = sd(NumBikes))
```

```
# A tibble: 2 x 3
  Holiday    mean    sd
  <fct>      <dbl> <dbl>
1 Holiday    529.  574.
2 No Holiday  739.  644.
```

```
sub_data |> # check rain by season
  group_by(Season) |>
  summarize(mean = mean(Rainfall), sd = sd(Rainfall))
```

```
# A tibble: 4 x 3
  Season    mean    sd
  <fct>    <dbl> <dbl>
1 Autumn  0.118  0.890
2 Spring  0.187  1.21
3 Summer  0.253  1.59
4 Winter  0.0328 0.423
```

```
sub_data |> # check snowfall by season
  group_by(Season) |>
  summarize(mean = mean(Snowfall), sd = sd(Snowfall))
```

```
# A tibble: 4 x 3
  Season    mean    sd
  <fct>    <dbl> <dbl>
1 Autumn  0.0635 0.522
2 Spring  0      0
3 Summer  0      0
4 Winter  0.248  0.698
```

We see that summer appears to have the highest number of bike rentals, and winter has the fewest. Bike rentals are more common on Non-Holidays, so they are likely being used to commute to work. Finally, we see from our tables that snowfall is most plentiful in winter, while rainfall is most plentiful in summer.

Next up, let's summarize across hours so that we can collapse each day into a single observation.

```
clean_data = sub_data |>
  group_by(Date, Season, Holiday) |>
  summarize(TotBikes = sum(NumBikes), TotRain = sum(Rainfall), TotSnow = sum(Snowfall), MeanTemp = mean(Temp), MeanHumid = mean(Humid))
```

`summarise()` has grouped output by 'Date', 'Season'. You can override using the `.groups` argument.

```
head(clean_data)
```

```
# A tibble: 6 x 12
# Groups:   Date, Season [6]
  Date      Season Holiday TotBikes TotRain TotSnow MeanTemp MeanHumid
  <date>    <fct>  <fct>    <int>   <dbl>   <dbl>    <dbl>    <dbl>
1 2017-12-01 Winter No Holiday   9539     0       0    -2.45     45.9
2 2017-12-02 Winter No Holiday   8523     0       0     1.32     62.0
3 2017-12-03 Winter No Holiday   7222     4       0     4.88     81.5
4 2017-12-04 Winter No Holiday   8729    0.1     0    -0.304     52.5
5 2017-12-05 Winter No Holiday   8307     0       0    -4.46     36.4
6 2017-12-06 Winter No Holiday   6669    1.3     8.6    0.0458     70.8
# i 4 more variables: MeanWindSpeed <dbl>, MeanVis <dbl>, MeanDewPoint <dbl>,
#   MeanSolar <dbl>
```

Great, now we have our final data set for training. Let's recreate our basic summaries and do some plots on this cleaned data. We'll also report some correlations. Let's first recreate our summary stats:

```
clean_data |> # check bike numbers by season
  group_by(Season) |>
  summarize(mean = mean(TotBikes), sd = sd(TotBikes))
```

```
# A tibble: 4 x 3
  Season  mean  sd
  <fct>   <dbl> <dbl>
1 Autumn 22099. 6711.
2 Spring 17910. 8357.
3 Summer 24818. 7297.
4 Winter  5413. 1808.
```

```
clean_data |> # check bike numbers by holiday
  group_by(Holiday) |>
  summarize(mean = mean(TotBikes), sd = sd(TotBikes))
```

```
# A tibble: 2 x 3
  Holiday      mean    sd
  <fct>      <dbl> <dbl>
1 Holiday  12700. 10504.
2 No Holiday 17727.  9862.
```

```
clean_data |> # check rain by season
  group_by(Season) |>
  summarize(mean = mean(TotRain), sd = sd(TotRain))
```

```
# A tibble: 4 x 3
  Season mean    sd
  <fct> <dbl> <dbl>
1 Autumn 2.81   8.61
2 Spring 4.49  12.7
3 Summer 6.08  17.0
4 Winter 0.788  3.28
```

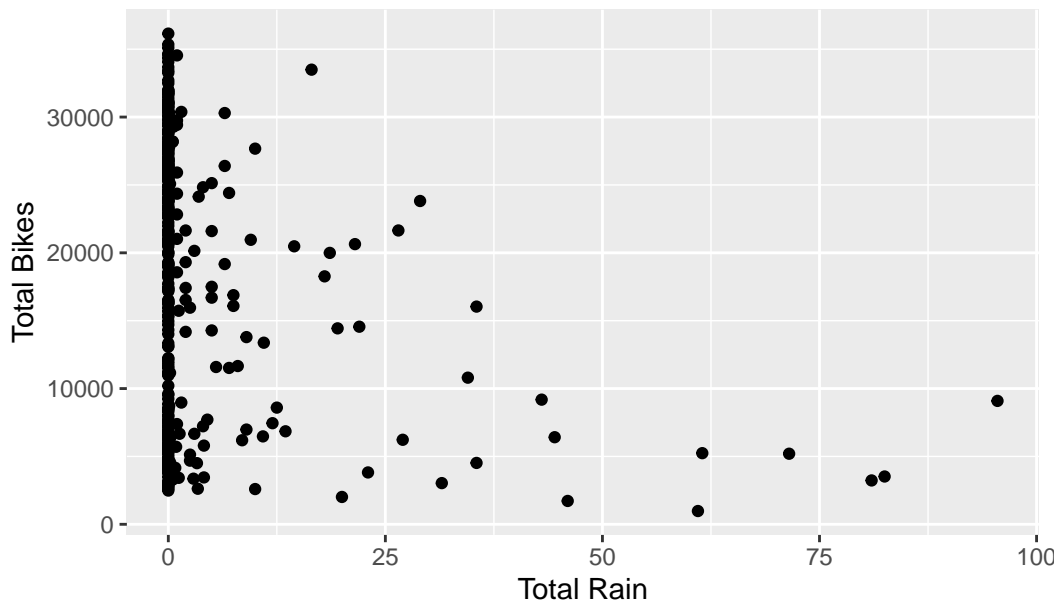
```
clean_data |> # check snowfall by season
  group_by(Season) |>
  summarize(mean = mean(TotSnow), sd = sd(TotSnow))
```

```
# A tibble: 4 x 3
  Season mean    sd
  <fct> <dbl> <dbl>
1 Autumn 1.52   9.83
2 Spring 0      0
3 Summer 0      0
4 Winter 5.94  14.0
```

We see that the same trends hold; more bikes in summer, more bikes on non-holidays, more rain in summer, and more snow in winter. Next, let's plot some of the numerical variables:

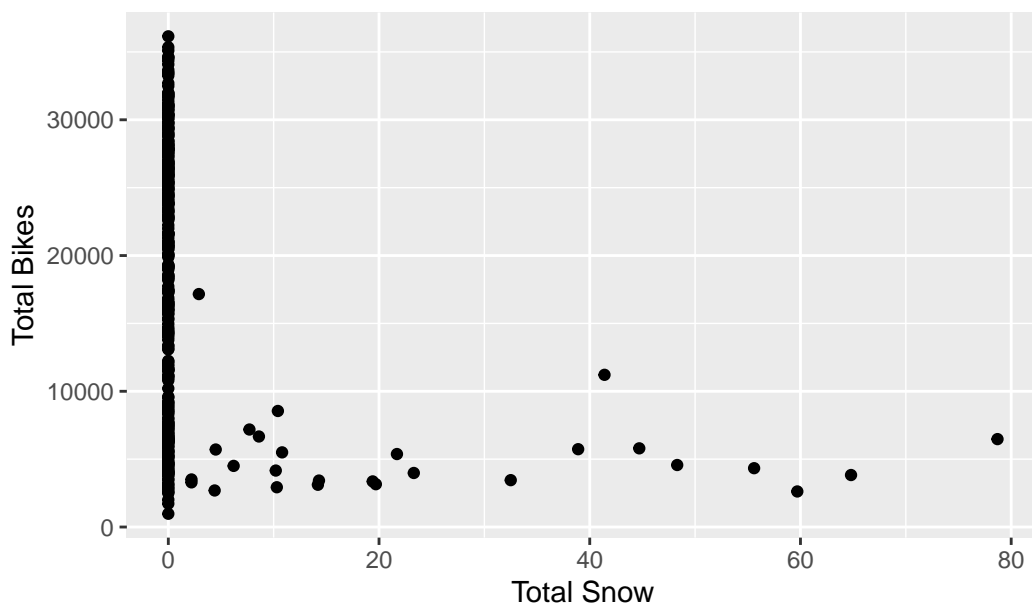
```
ggplot(clean_data, aes(x = TotRain, y = TotBikes)) +
  geom_point() +
  labs(x = 'Total Rain', y = 'Total Bikes', title = 'Bike Rentals by Amount of rain')
```


Bike Rentals by Amount of rain



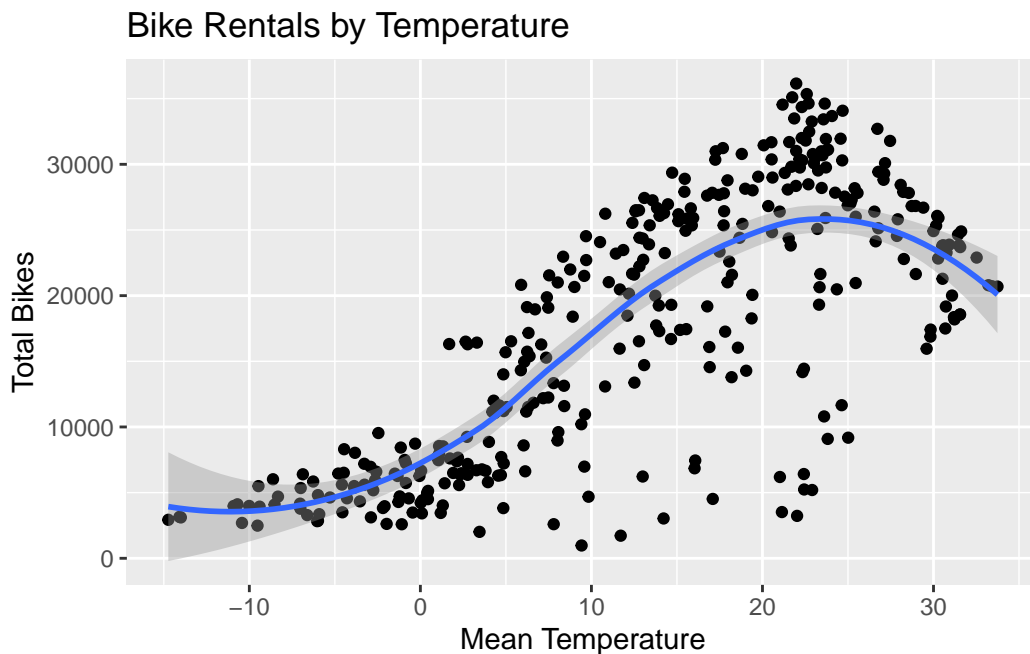
```
ggplot(clean_data, aes(x = TotSnow, y = TotBikes)) +  
  geom_point() +  
  labs(x = 'Total Snow', y = 'Total Bikes', title = 'Bike Rentals by Amount of snow')
```

Bike Rentals by Amount of snow



```
ggplot(clean_data, aes(x = MeanTemp, y = TotBikes)) +
  geom_point() +
  geom_smooth() +
  labs(x = 'Mean Temperature', y = 'Total Bikes', title = 'Bike Rentals by Temperature')
```

`geom_smooth()` using method = 'loess' and formula = 'y ~ x'



These plots make sense. We see a lot more bike rentals when there is low snow and low rain. Also, we see that the number of bikes increase as temperature increases, until a certain point. These plots appear to be accurately capturing info about how weather influence bike riding behavior!

Finally, let's plot some correlations:

```
cor(clean_data[sapply(clean_data, is.numeric)])
```

	TotBikes	TotRain	TotSnow	MeanTemp	MeanHumid
TotBikes	1.00000000	-0.23910905	-0.26529110	0.753076732	0.03588697
TotRain	-0.23910905	1.00000000	-0.02313404	0.144517274	0.52864263
TotSnow	-0.26529110	-0.02313404	1.00000000	-0.266963662	0.06539191
MeanTemp	0.75307673	0.14451727	-0.26696366	1.00000000	0.40416749
MeanHumid	0.03588697	0.52864263	0.06539191	0.404167486	1.00000000

MeanWindSpeed	-0.19288142	-0.10167578	0.02088156	-0.260721792	-0.23425778
MeanVis	0.16599375	-0.22199387	-0.10188902	0.002336683	-0.55917733
MeanDewPoint	0.65047655	0.26456621	-0.20955286	0.962796255	0.63204729
MeanSolar	0.73589290	-0.32270413	-0.23343056	0.550274301	-0.27444967
	MeanWindSpeed	MeanVis	MeanDewPoint	MeanSolar	
TotBikes	-0.19288142	0.165993749	0.6504765	0.73589290	
TotRain	-0.10167578	-0.221993866	0.2645662	-0.32270413	
TotSnow	0.02088156	-0.101889019	-0.2095529	-0.23343056	
MeanTemp	-0.26072179	0.002336683	0.9627963	0.55027430	
MeanHumid	-0.23425778	-0.559177334	0.6320473	-0.27444967	
MeanWindSpeed	1.00000000	0.206022636	-0.2877032	0.09612635	
MeanVis	0.20602264	1.000000000	-0.1535516	0.27139591	
MeanDewPoint	-0.28770322	-0.153551591	1.0000000	0.38315713	
MeanSolar	0.09612635	0.271395906	0.3831571	1.00000000	

We can see some impressive correlations here. For example, Mean Temp, Mean Dew Point, and Mean Solar are all quite positively correlated with total number of bike rentals. These are likely all tied into the weather we commented on above - when the weather is nice, people are more likely to rent bikes, thus leading to higher correlation values! We should be able to do some nice prediction on this data given the strengths of these correlations, although we might be at risk of overfitting if we aren't careful.

Modeling the Data